

Assignment 2

NAME: Nicholas Tolley

DUE DATE: February 28th, 6pm

Problem 1 (100 pts)

In the earnings dataset you can find salary (*earn*) and some socio-demographic characteristics of each subject, including variables such as *height*, *weight*, gender (*male*), *ethnicity*, *education*, mother's (*mother_education*) and father's education (*father_education*), *walk* (e.g. walking time), *exercise*, if they smoke or not (*smokenow*), *tense*, *angry* and *age*.

The dataset can be found in Canvas in the Data folder (file name: earnings.csv):

- (a) (10 points) Subset the data and consider only the variables: *education*, *mother_education*, *father_education*, *walk*, *exercise*, *tense*, *angry*, *weight*, *height*. Check the correlation by performing a figure similar to Figure 1 below (make sure not to use the default colours but rather choose your own). Take special care to the labels and legend. What can you say about the results? What would you expect from a linear regression model (hint: there are some variables to be excluded/included in the model)? Perform a test statistic for the correlation between *earn* and *education*, write the hypothesis test and the results you will obtain.

```
library(ggplot2)
library(boot)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v tibble 3.1.8      v dplyr 1.1.0
## v tidyr 1.3.0      v stringr 1.5.0
## v readr 2.1.3      v forcats 1.0.0
## v purrr 1.0.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()

library(GGally)

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg ggplot2

library(tidyr)
library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##   select
```

```
library(caret)

## Loading required package: lattice
##
## Attaching package: 'lattice'
##
## The following object is masked from 'package:boot':
##
##     melanoma
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift

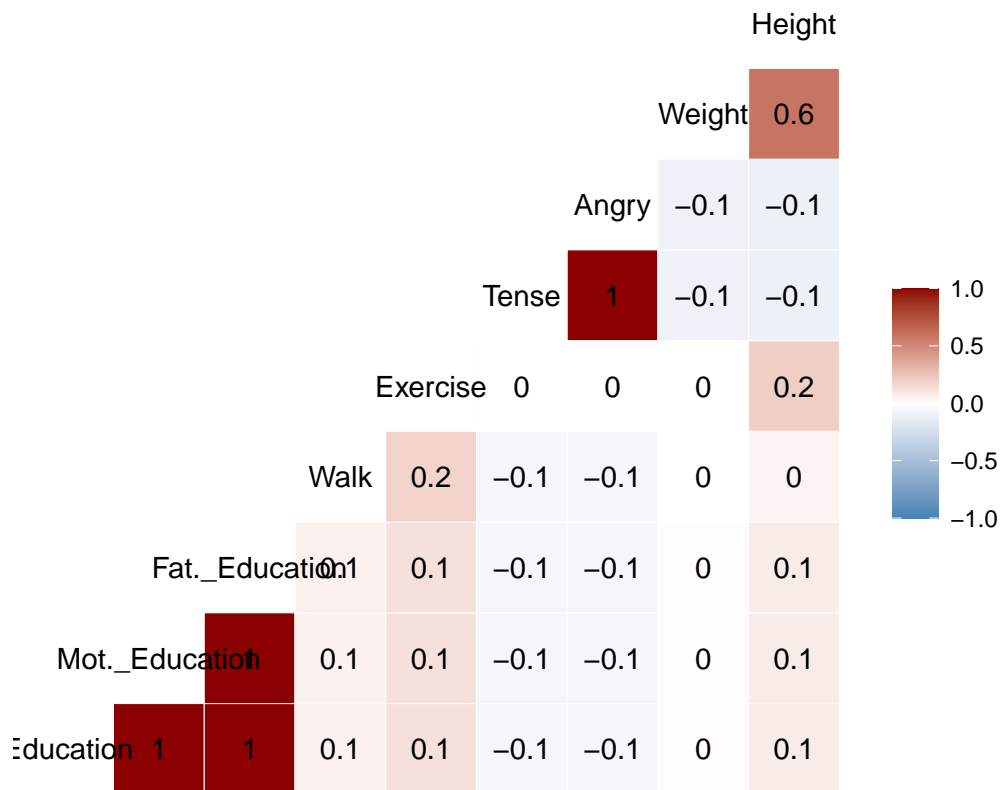
df <- read.csv('earnings.csv')
```

The code below stores a subset of the dataset with the columns indicated above. Since we are calculating the correlation between columns, rows with missing values in any column are removed.

```
subset_cols <- c('education', 'mother_education', 'father_education',
                 'walk', 'exercise', 'tense', 'angry', 'weight', 'height')
label_names <- c('Education', 'Mot._Education', 'Fat._Education', 'Walk', 'Exercise', 'Tense',
                 'Angry', 'Weight', 'Height')
df_subset <- subset(df, select=subset_cols)
df_subset <- drop_na(df_subset)
colnames(df_subset) <- label_names
```

Next we can calculate the correlation between the columns contained in the subset, and visualize the result as a heatmap.

```
corr_matrix <- cor(df_subset)
ggcorr(df_subset, label=TRUE, low="steelblue", mid="white", high="darkred")
```



As we can see there are several variables that exhibit a near perfect correlation with one another. The most highly correlated columns include:

education <-> mother_education

education <-> father_education

father_education <-> mother_education

tense <-> angry

If we were to build a linear regression model, we would need to remove 2 of the education variables, and either the tense or angry variable. This is because the dataset exhibits what is known as multicollinearity, in other words there is redundant information in the columns. If we were to try and create a linear model on the full dataset, there would not be a unique combination of regression (“beta”) coefficients that minimize the residual error. For example, the same coefficient could be assigned to the angry or tense columns.

The code below calculates the correlation coefficient, and associated p-value, between the earn and education columns. The p-value refers to the probability of the null hypothesis that these two variables are uncorrelated (correlation=0).

```
cor.test(df$earn, df$education)

##
## Pearson's product-moment correlation
##
## data: df$earn and df$education
## t = 13.748, df = 1812, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
```

```
## 0.2650564 0.3484265
## sample estimates:
##      cor
## 0.3073311
```

We can see that despite the correlation being relatively low ($\text{cor}=0.3$), the result is highly significant with a p-value $< 2.2\text{e-}16$. We can therefore reject the null hypothesis that the “earn” and “education” variables are uncorrelated with one another.



Figure 1: Correlation

- (b) (10 points) Perform a linear regression model using the variable *earn* as the dependent variable and years of education *education* as the independent variable. What can you say about this covariate? Is it significant? Write down the hypothesis test. Plot the linear regression you have obtained in ggplot by using a subset of the data. This subset is obtained by restricting the variable *earn* to be less than $2\text{e}+05$ (similar to Figure 2 below)

The code below creates a linear model that predicts “earn” by the education covariate. From the linear model there are three potential hypothesis tests with the following null hypotheses:

- The intercept of the linear model is zero
- The beta coefficient of the linear model is zero

- The model has the same explanatory power (residual error) as constant (flat line) model

```
earn_fit <- lm(df$earn ~ df$education)
summary(earn_fit)
```

```
##
## Call:
## lm(formula = df$earn ~ df$education)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -34051 -12373  -3212   7207 382207
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -14724.1      2657.0  -5.542 3.43e-08 ***
## df$education  2709.7        197.1  13.748 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21460 on 1812 degrees of freedom
## (2 observations deleted due to missingness)
## Multiple R-squared:  0.09445,    Adjusted R-squared:  0.09395
## F-statistic:   189 on 1 and 1812 DF,  p-value: < 2.2e-16
```

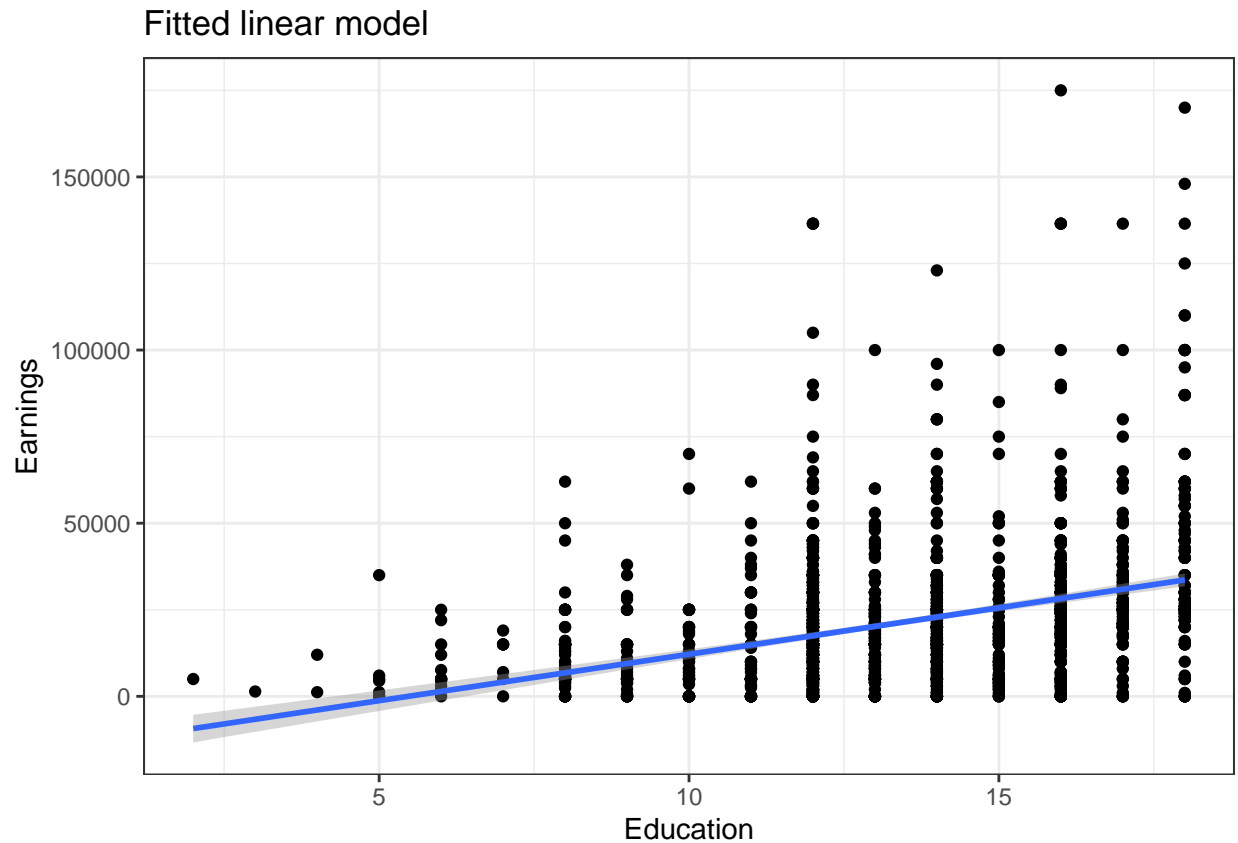
As we can see from the summary of the model output above, both the intercept and the slope are highly significant indicating that the *earn* variable is predicted well by the *education* variable.

Next using the subset of the with *earn* < 2e+05, we can visualize how well our linear model explains *earn* using just the *education* covariate

```
earn_subset <- subset(df, df$earn < 2e5)
```

```
ggplot(earn_subset, aes(x=education, y=earn)) + geom_point(na.rm=TRUE) + geom_smooth(method='lm', na.rm=TRUE)
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



- (c) (20 points) Draw the qqplot by using the library ggplot for the model obtained in point b. Then perform the qqplot (using the library ggplot) for the two different groups of sex (similar to Figure 3 below). Take special care to the legend and the label. What can you say about this plot?

```
model_residuals <- resid(earn_fit)
p <- seq(0.05, 0.95, 0.01)
sample_quantiles <- quantile(model_residuals, p)
theoretical_quantiles <- qnorm(p, mean=mean(model_residuals), sd=sd(model_residuals))
qqplot(theoretical_quantiles, sample_quantiles) + geom_abline() + labs(title='Residuals Q-Q Plot', y='Residuals')

## Warning: `qqplot()` was deprecated in ggplot2 3.4.0.
```


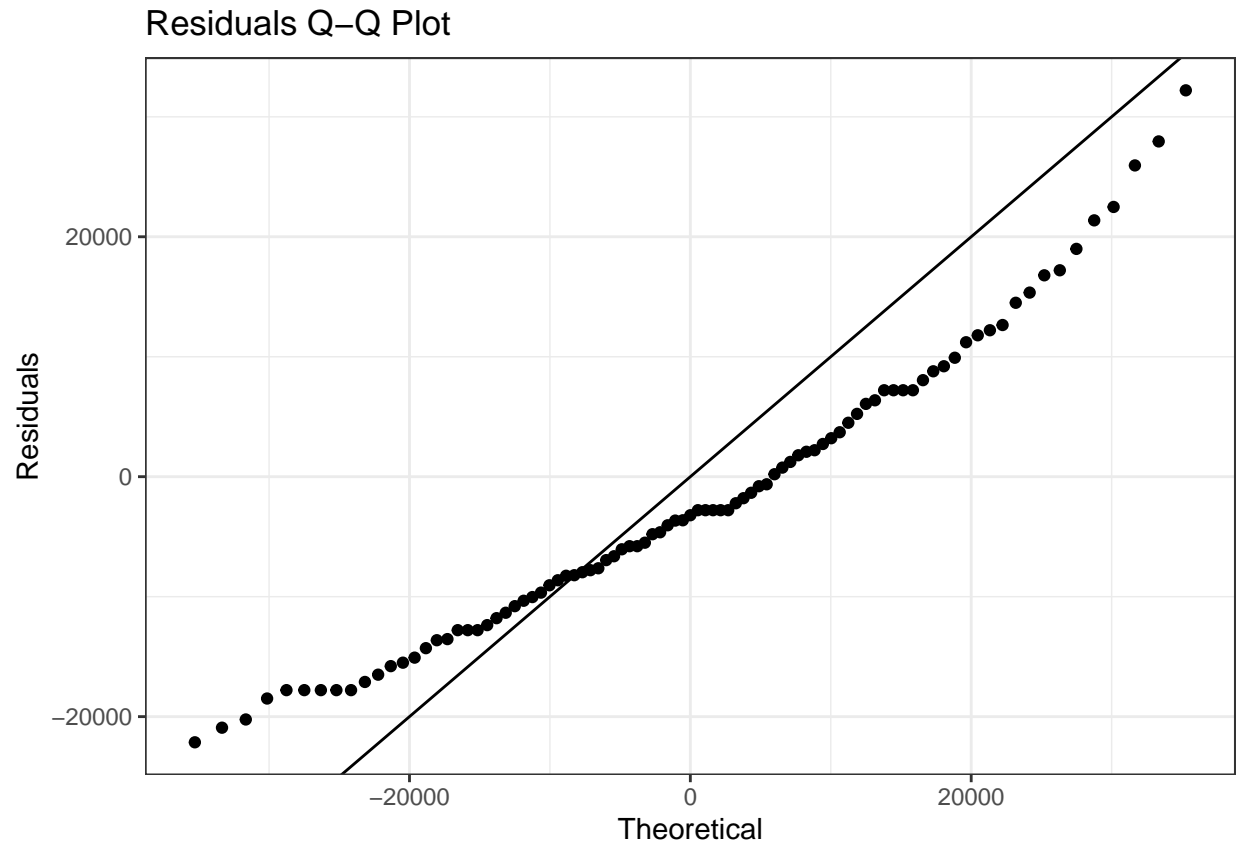


Figure 2: Linear Regression



```
earn_quantiles_male <- quantile(subset(df, df$male==1)$earn, p, na.rm=TRUE)
earn_quantiles_female <- quantile(subset(df, df$male==0)$earn, p, na.rm=TRUE)

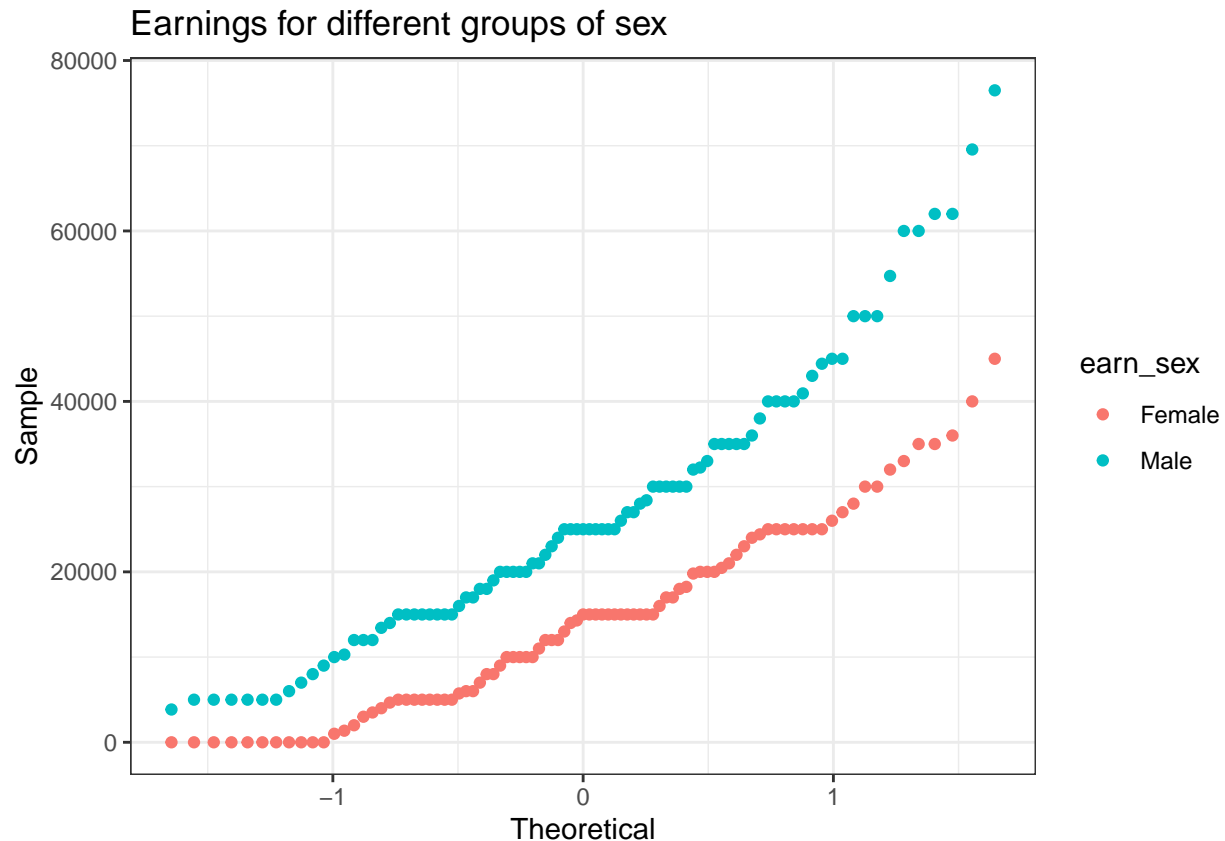
theoretical_quantiles_male <- qnorm(p, mean=mean(earn_quantiles_male),
                                   sd=sd(earn_quantiles_male))
theoretical_quantiles_female <- qnorm(p, mean=mean(earn_quantiles_female),
                                     sd=sd(earn_quantiles_female))

theoretical_quantiles_male <- qnorm(p, mean=0, sd=1)
theoretical_quantiles_female <- qnorm(p, mean=0, sd=1)

earn_quantiles <- c(earn_quantiles_male, earn_quantiles_female)
theoretical_quantiles <- c(theoretical_quantiles_male, theoretical_quantiles_female)
earn_sex <- c(rep('Male', each=length(earn_quantiles_male)),
             rep('Female', each=length(earn_quantiles_female)))

earn_sex_df <- data.frame(earn_quantiles, theoretical_quantiles, earn_sex)

ggplot(earn_sex_df, aes(x=theoretical_quantiles, y=earn_quantiles, color=earn_sex)) +
  geom_point() + labs(title='Earnings for different groups of sex', y='Sample', x='Theoretical') +
  theme_bw()
```

- (d) (20 points) Perform in R the backward and forward procedure to select the covariates, remember to remove the rows with missing values. Did you obtain the same or different results from the two different procedures, please explain. Which procedure would you prefer? Comment what you discovered and the theoretical implications. Just for the backward solution compute the RSS and show the trend of RSS for beta1 in a plot by using ggplot in R (similar to Figure 4). (Hint: For RSS plot, set the range of x-axis to be [0,1000]).

The code below performs the backward procedure to select...

```
fit1 <- lm(earn~., data=drop_na(df))
fit2 <- lm(earn~1, data=drop_na(df))
mod_backward <- stepAIC(fit1, direction="backward", scope=list(upper=fit1, lower=fit2))
```

```
## Start: AIC=28708.32
## earn ~ height + weight + male + ethnicity + education + mother_education +
##       father_education + walk + exercise + smokenow + tense + angry +
##       age
##
##
## Step: AIC=28708.32
## earn ~ height + weight + male + ethnicity + education + mother_education +
##       father_education + walk + exercise + smokenow + tense + age
##
##
## Step: AIC=28708.32
## earn ~ height + weight + male + ethnicity + education + mother_education +
##       walk + exercise + smokenow + tense + age
```




Figure 3: QQplot for different groups

```

##
##
## Step: AIC=28708.32
## earn ~ height + weight + male + ethnicity + education + walk +
## exercise + smokenow + tense + age
##
##          Df Sum of Sq      RSS   AIC
## - ethnicity 3 1.4940e+09 6.4531e+11 28706
## - walk      1 3.1861e+07 6.4385e+11 28706
## - weight    1 5.1199e+07 6.4387e+11 28706
## - smokenow  1 2.5692e+08 6.4407e+11 28707
## - exercise  1 3.2969e+08 6.4415e+11 28707
## <none>              6.4382e+11 28708
## - tense     1 1.4955e+09 6.4531e+11 28710
## - height    1 2.2805e+09 6.4610e+11 28711
## - age       1 2.0675e+10 6.6449e+11 28752
## - male      1 2.5472e+10 6.6929e+11 28762
## - education 1 6.6279e+10 7.1009e+11 28847
##
## Step: AIC=28705.66
## earn ~ height + weight + male + education + walk + exercise +
## smokenow + tense + age
##
##          Df Sum of Sq      RSS   AIC
## - walk      1 6.3440e+07 6.4537e+11 28704
## - weight    1 7.7849e+07 6.4539e+11 28704
## - smokenow  1 3.0394e+08 6.4561e+11 28704
## - exercise  1 3.4289e+08 6.4565e+11 28704
## <none>              6.4531e+11 28706
## - tense     1 1.4071e+09 6.4672e+11 28707
## - height    1 2.8270e+09 6.4814e+11 28710
## - age       1 2.2170e+10 6.6748e+11 28752
## - male      1 2.4948e+10 6.7026e+11 28758
## - education 1 6.7532e+10 7.1284e+11 28847
##
## Step: AIC=28703.8
## earn ~ height + weight + male + education + exercise + smokenow +
## tense + age
##
##          Df Sum of Sq      RSS   AIC
## - weight    1 7.4075e+07 6.4545e+11 28702
## - exercise  1 3.0181e+08 6.4568e+11 28702
## - smokenow  1 3.1560e+08 6.4569e+11 28702
## <none>              6.4537e+11 28704
## - tense     1 1.4239e+09 6.4680e+11 28705
## - height    1 2.8131e+09 6.4819e+11 28708
## - age       1 2.2147e+10 6.6752e+11 28750
## - male      1 2.4931e+10 6.7030e+11 28756
## - education 1 6.7475e+10 7.1285e+11 28845
##
## Step: AIC=28701.97
## earn ~ height + male + education + exercise + smokenow + tense +
## age
##

```

```
##           Df Sum of Sq      RSS      AIC
## - smokenow  1 3.3784e+08 6.4579e+11 28701
## - exercise  1 3.4429e+08 6.4579e+11 28701
## <none>                6.4545e+11 28702
## - tense      1 1.4060e+09 6.4685e+11 28703
## - height     1 2.9325e+09 6.4838e+11 28706
## - age        1 2.2336e+10 6.6778e+11 28749
## - male       1 2.5233e+10 6.7068e+11 28755
## - education  1 6.7691e+10 7.1314e+11 28844
##
## Step: AIC=28700.72
## earn ~ height + male + education + exercise + tense + age
##
##           Df Sum of Sq      RSS      AIC
## - exercise  1 2.8870e+08 6.4607e+11 28699
## <none>                6.4579e+11 28701
## - tense      1 1.4833e+09 6.4727e+11 28702
## - height     1 2.9449e+09 6.4873e+11 28705
## - age        1 2.2071e+10 6.6786e+11 28747
## - male       1 2.5335e+10 6.7112e+11 28754
## - education  1 6.7697e+10 7.1348e+11 28842
##
## Step: AIC=28699.37
## earn ~ height + male + education + tense + age
##
##           Df Sum of Sq      RSS      AIC
## <none>                6.4607e+11 28699
## - tense      1 1.4303e+09 6.4750e+11 28700
## - height     1 3.0014e+09 6.4908e+11 28704
## - age        1 2.2443e+10 6.6852e+11 28746
## - male       1 2.6396e+10 6.7247e+11 28755
## - education  1 6.9666e+10 7.1574e+11 28845
```

Likewise, we can perform the forward pass by...

```
mod_forward <- stepAIC(fit2, direction="forward", scope=list(upper=fit1, lower=fit2))
```

```
## Start: AIC=29031.73
## earn ~ 1
##
##           Df Sum of Sq      RSS      AIC
## + male      1 7.9739e+10 7.3974e+11 28886
## + education  1 7.4527e+10 7.4495e+11 28896
## + mother_education 1 7.4527e+10 7.4495e+11 28896
## + father_education 1 7.4527e+10 7.4495e+11 28896
## + height     1 5.9505e+10 7.5997e+11 28925
## + weight     1 2.9325e+10 7.9016e+11 28981
## + age        1 1.3370e+10 8.0611e+11 29010
## + exercise   1 5.0787e+09 8.1440e+11 29025
## + tense      1 3.9410e+09 8.1554e+11 29027
## + angry      1 3.9410e+09 8.1554e+11 29027
## + ethnicity  3 4.8139e+09 8.1467e+11 29029
## <none>                8.1948e+11 29032
## + smokenow   1 9.1854e+08 8.1856e+11 29032
## + walk       1 9.0848e+08 8.1857e+11 29032
```

```

##
## Step: AIC=28886.32
## earn ~ male
##
##           Df Sum of Sq      RSS      AIC
## + education      1 6.9797e+10 6.6994e+11 28746
## + mother_education 1 6.9797e+10 6.6994e+11 28746
## + father_education 1 6.9797e+10 6.6994e+11 28746
## + age            1 1.7756e+10 7.2198e+11 28853
## + height         1 3.6894e+09 7.3605e+11 28881
## + ethnicity      3 5.3068e+09 7.3443e+11 28882
## <none>                7.3974e+11 28886
## + weight         1 9.4776e+08 7.3879e+11 28886
## + smokenow       1 9.1854e+08 7.3882e+11 28886
## + tense          1 3.4937e+08 7.3939e+11 28888
## + angry          1 3.4937e+08 7.3939e+11 28888
## + walk           1 2.3203e+08 7.3951e+11 28888
## + exercise       1 9.0739e+07 7.3965e+11 28888
##
## Step: AIC=28745.61
## earn ~ male + education
##
##           Df Sum of Sq      RSS      AIC
## + age            1 1.9518e+10 6.5043e+11 28705
## + ethnicity      3 3.5226e+09 6.6642e+11 28744
## + height         1 1.4225e+09 6.6852e+11 28744
## + weight         1 1.0396e+09 6.6890e+11 28745
## <none>                6.6994e+11 28746
## + exercise       1 5.4813e+08 6.6940e+11 28746
## + smokenow       1 1.1241e+08 6.6983e+11 28747
## + walk           1 3.1565e+06 6.6994e+11 28748
## + tense          1 2.5823e+06 6.6994e+11 28748
## + angry          1 2.5823e+06 6.6994e+11 28748
##
## Step: AIC=28705.03
## earn ~ male + education + age
##
##           Df Sum of Sq      RSS      AIC
## + height         1 2921837319 6.4750e+11 28700
## + tense          1 1350718589 6.4908e+11 28704
## + angry          1 1350718589 6.4908e+11 28704
## <none>                6.5043e+11 28705
## + smokenow       1 362030699 6.5006e+11 28706
## + ethnicity      3 2093480782 6.4833e+11 28706
## + exercise       1 287699226 6.5014e+11 28706
## + weight         1 137467965 6.5029e+11 28707
## + walk           1 33481537 6.5039e+11 28707
##
## Step: AIC=28700.55
## earn ~ male + education + age + height
##
##           Df Sum of Sq      RSS      AIC
## + tense          1 1430271509 6.4607e+11 28699
## + angry          1 1430271509 6.4607e+11 28699

```

```
## <none> 6.4750e+11 28700
## + smokenow 1 357020068 6.4715e+11 28702
## + exercise 1 235648665 6.4727e+11 28702
## + weight 1 111208434 6.4739e+11 28702
## + walk 1 41917712 6.4746e+11 28702
## + ethnicity 3 1521906378 6.4598e+11 28703
##
## Step: AIC=28699.37
## earn ~ male + education + age + height + tense
##
## Df Sum of Sq RSS AIC
## <none> 6.4607e+11 28699
## + exercise 1 288700139 6.4579e+11 28701
## + smokenow 1 282249180 6.4579e+11 28701
## + weight 1 137259702 6.4594e+11 28701
## + walk 1 25908265 6.4605e+11 28701
## + ethnicity 3 1607572422 6.4447e+11 28702
```

To find the RSS associated with changing the beta 1 coefficient, we will first create a function that manually performs the matrix multiplication using the weights arrived at from the backward pass, and then updates the coefficient based on our input.


```
beta_update <- function(new_val, mod, d){
  beta_coef <- mod[["coefficients"]]
  beta_coef[2] <- new_val
  mod_d <- subset(d, select=names(mod[["coefficients"]][c(2,3,4,5,6)]))
  mod_d <- data.frame(intercept=rep(1, nrow(mod_d)), mod_d)
  pred <- as.matrix(mod_d) %*% beta_coef

  res <- sum((d$earn - pred)^2)
  return(res)
}

rss_out <- rep(0, 1000)
for (i in 1:1000){
  rss_out[i] <- beta_update(i, mod_backward, drop_na(df))
}

rss_df <- data.frame(RSS=rss_out, beta1=c(1:1000))

ggplot(data=rss_df, aes(x=beta1, y=rss_out)) + geom_point() +
  labs(title="RSS with altered height coefficient", y="RSS") +
  geom_vline(xintercept=mod_backward[["coefficients"]][2], color="red")
```



- (e) (20 points) Perform a bootstrap of 500 samples for beta 1 (*height*), beta 2 (*male*), and beta 3 (*education*) for the coefficient obtained in the backward procedure in point d. Plot the beta coefficients that you have obtained with histograms with ggplot (similar to Figure 5). Remember to use the data without missing values.

```
fc <- function(d, i){
  d2 <- d[i,]
  boot_fit <- lm(earn ~ height + male + education + tense + age, data=d2)

  return(boot_fit[['coefficients']])
}

earn_boot <- boot(drop_na(df), fc, R=500)

earn_boot_coef <- data.frame(earn_boot[["t"]])
colnames(earn_boot_coef) <- names(earn_boot[["t0"]])
ggplot(earn_boot_coef, aes(x=height)) + geom_histogram(bins=50) +
  labs(title='Height coefficient bootstrap', x='Beta Coefficient', y='Count') +
  geom_vline(xintercept=earn_boot[["t0"]][["height"]], color="red") +
  theme_bw()
```


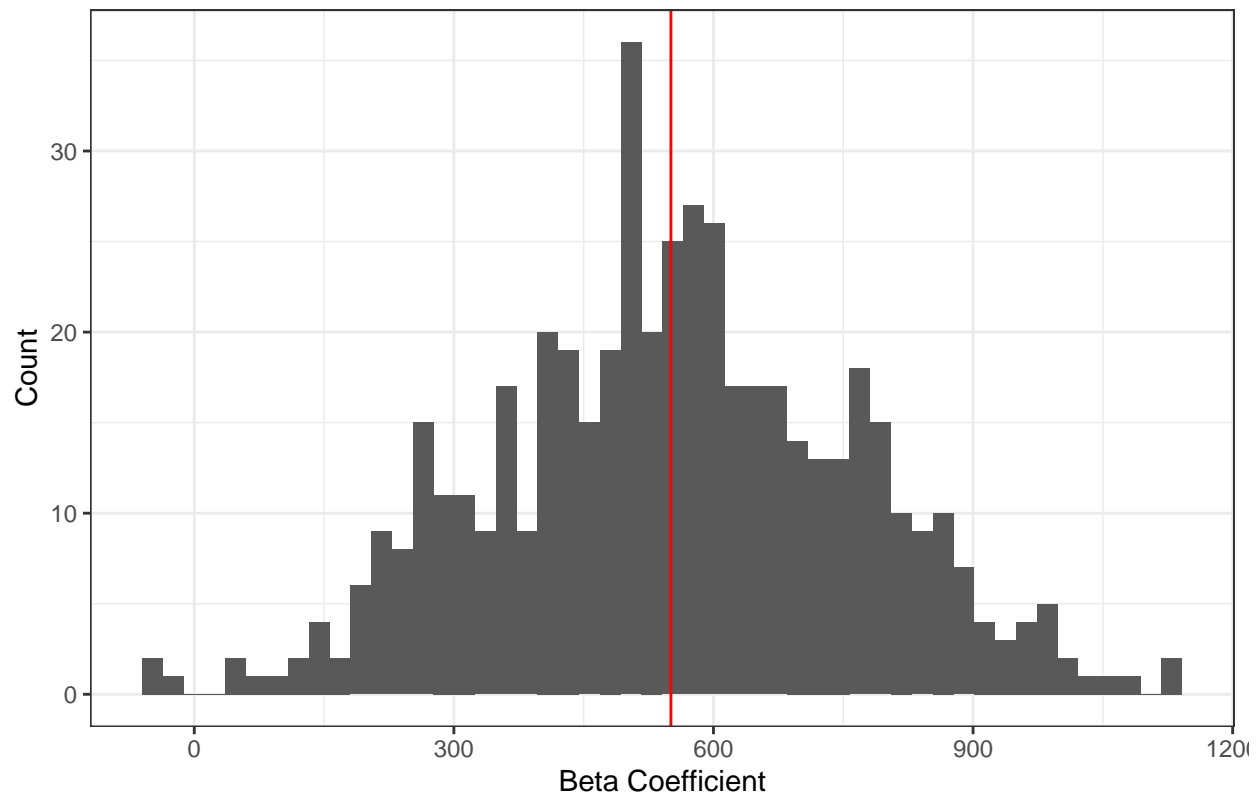



Figure 4: RSS for the backward procedure

Height coefficient bootstrap



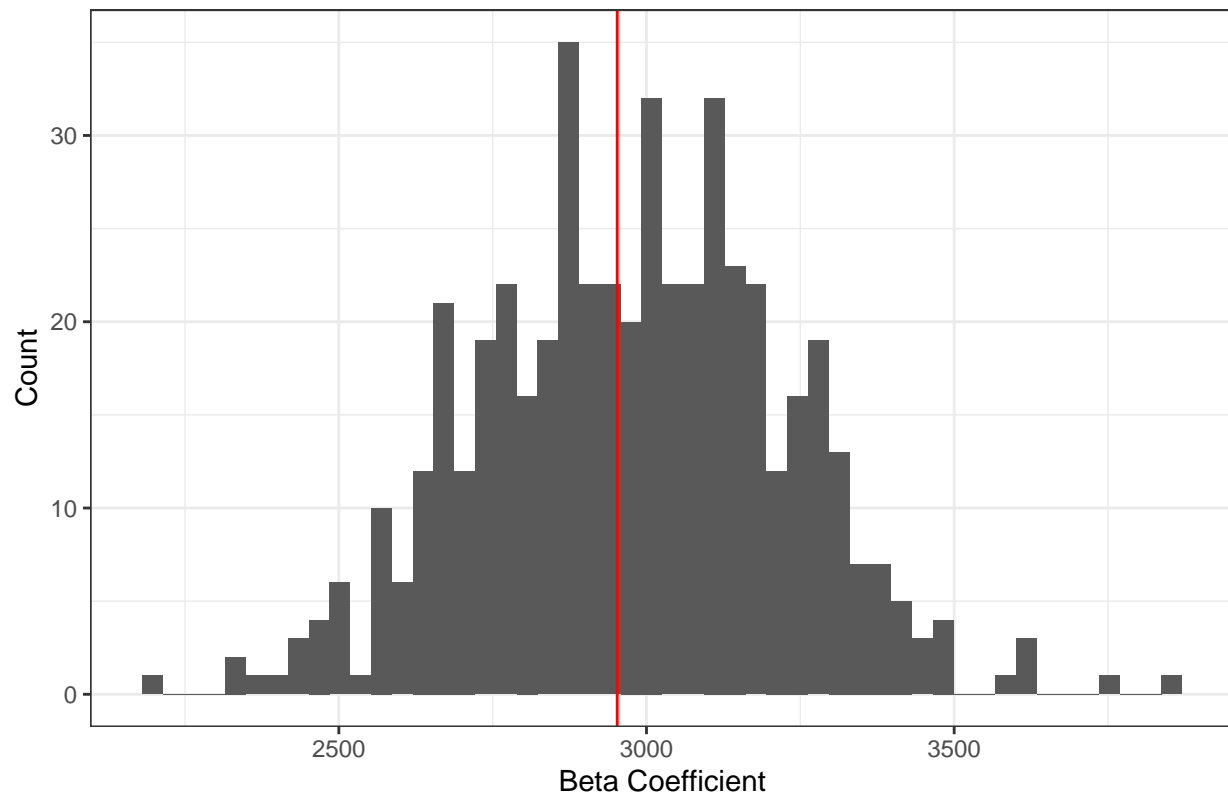
```
earn_boot_coef <- data.frame(earn_boot[["t"]])
colnames(earn_boot_coef) <- names(earn_boot[["t0"]])
ggplot(earn_boot_coef, aes(x=male)) + geom_histogram(bins=50) +
  labs(title='Male coefficient bootstrap', x='Beta Coefficient', y='Count') +
  geom_vline(xintercept=earn_boot[["t0"]][["male"]], color="red") +
  theme_bw()
```

Male coefficient bootstrap



```
earn_boot_coef <- data.frame(earn_boot[["t"]])
colnames(earn_boot_coef) <- names(earn_boot[["t0"]])
ggplot(earn_boot_coef, aes(x=education)) + geom_histogram(bins=50) +
  labs(title='Education coefficient bootstrap', x='Beta Coefficient', y='Count') +
  geom_vline(xintercept=earn_boot[["t0"]][["education"]], color="red") +
  theme_bw()
```

Education coefficient bootstrap



- (f) (20 points) Compute the LOO and K-fold cross validation and write the results. Compute the mean square error for both the LOO and the K-fold cross validation. Then plot the prediction against the true value for LOO, using ggplot. Describe the results. Remember to use the data without missing values.

```
train.control <- trainControl(method = "LOOCV")
model_loocv <- train(earn~height + male + education, data=drop_na(df),
                     method="lm", trControl=train.control)
print(model_loocv)
```

```
## Linear Regression
##
## 1440 samples
##    3 predictor
##
## No pre-processing
## Resampling: Leave-One-Out Cross-Validation
## Summary of sample sizes: 1439, 1439, 1439, 1439, 1439, 1439, ...
## Resampling results:
##
##    RMSE      Rsquared    MAE
## 21612.29  0.1792554 13070.81
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
train.control <- trainControl(method = "cv", number=10)
model_cv <- train(earn~height + male + education, data=drop_na(df),
```




Figure 5: Bootstrap Results

```

      method="lm", trControl=train.control)
print(model_cv)

## Linear Regression
##
## 1440 samples
##    3 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1295, 1297, 1294, 1298, 1296, 1296, ...
## Resampling results:
##
##    RMSE      Rsquared   MAE
## 20682.8  0.2035937 13063.7
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
pred <- predict(model_cv, newdata = drop_na(df))
cv_df <- data.frame(true_earn=drop_na(df)$earn, pred_earn=pred)
ggplot(cv_df, aes(x=true_earn, y=pred_earn)) + geom_point() +
  labs(title='Linear Model True Earnings vs Predicted Earnings', x='True', y='Predicted') + theme_bw()

```

